

# Recognition Class 12



# Kalman Filter

- Weighted sum:  $\hat{x}_t = \hat{x}_t^- + K_t(z_t - H\hat{x}_t^-)$
- Questions:



- How to set the weight between prediction and observations?

$$k_G = \frac{E_{EST}}{E_{EST} + E_{MEA}}$$

- How to set the variance?

$$E_{EST}(t) = (1 - k_G)E_{EST}(t-1)$$

- How to predict ?



# Simple Example

- Known system's **linear** dynamic model:

- E.g.,  $x_t = \begin{pmatrix} p_t \\ \dot{p}_t \end{pmatrix},$

$\dot{p}$  is velocity

$$\dot{p}_{t+1} = \dot{p}_t \text{ and } p_{t+1} = p_t + \Delta T \dot{p}_t + w_t$$

- Known linear mapping between the state  $x_t$  and the observation,  $z_t$ :

- E.g.,  $z_t = p_t + v_t$

- Measurement and estimation errors:

- $v_t$  and  $w_t$  are Gaussian noise

# Assumptions

- Known system's **linear** dynamic model, with white noise  $\sim G(0, Q)$

$$x_t = Ax_{t-1} + Bu_t + w_{t-1}$$

- $A$  is  $n \times n$ ,  $B$  is  $n \times \ell$
- Known **linear** transformation of the states to the measurements with white noise  $\sim G(0, R)$ :

$$z_t = Hx_t + v_t$$

- $H$  is an  $m \times n$  matrix

$A, B, H, Q, R$  are assumed to be known

# Goal

- Compute  $x_t$
- Given:
  - Initial state
  - Previous measurements
  - Known (or learned) linear models:  $A, B, H$
- Minimize the error (its covariance)  
between the correct and the computed  $x_t$

$$\begin{aligned}x_t &= Ax_{t-1} + Bu_t + w_{t-1}, & w &\in G(0, Q) \\z_t &= Hx_t + v_t, & v_t &\in G(0, R)\end{aligned}$$



# Notation

- A priori:
  - state estimation:  $\hat{x}_t^-$
  - error:  $e_t^- = x_t - \hat{x}_t^-$
  - covariance:  $\Sigma_t^- = E(e_t^- e_t^{-T})$
- A posteriori, given  $z_t$  :
  - state estimation:  $\hat{x}_t$
  - error:  $e_t = (x_t - \hat{x}_t)$
  - covariance:  $\Sigma_t = E(e_t e_t^T)$

We would like to minimize it

# The Discrete Kalman Filter



## Predict

(a priori estimate)

1. Predict the state ahead:

$$\hat{x}_t^- = A\hat{x}_{t-1}^- + Bu$$

2. Predict the error covariance ahead:

$$\Sigma_t^- = A \Sigma_{t-1} A^T + Q$$

## Update

(a posteriori estimate)

1. Kalman gain  $K_t$  is:

$$k_t = \frac{E_{est}}{E_{est} + E_{mea}}$$

$$K_t = \Sigma_t^- H^T (H \Sigma_t^- H^T + R)^{-1}$$

2. Update the state estimate:

$$\hat{x}_t = \hat{x}_t^- + K_t(z_t - H\hat{x}_t^-)$$

3. Update the error covariance:

$$\Sigma_t = (I - K_t H) \Sigma_t^-$$

$$x_t = Ax_{t-1} + Bu_t + w_{t-1}, \quad w \in G(0, Q)$$

$$z_t = Hx_t + v_t, \quad v_t \in G(0, R)$$

# Up to Here



- Up to here: only brief outline of Kalman filter
- Many extensions exists
- Detailed:  
<https://www.kalmanfilter.net/default.aspx>
- See more references at the of the presentation



# Non-Parametric Prediction


- Motivation: limitations of Kalman filter
- What are they?

# Particle Filters




- Can represent distribution with set of weighted samples (“*particles*”)
- Allows us to maintain **multiple hypotheses**

# Particle Filters



# Contour Prediction

- Object model
- Object shape
- Video



# More on Tracking

- Tracklets
- Evaluation



# Recognition



# What is Recognition

## Learning:

See one or more images of a given object: Building a model

## Recognition:

Recognize the object in a **novel** image:  
Determine the object (model) in the image

# Simple Correlation

The model:


- An image of the object
- A set of images of the object

**Recognition:** compare a new image to each of the models' images

- Use nearest neighbor
- Use k-nearest neighbors




# Identifiaction: Simple Correlation



A chair




Find the chair





Output of normalized  
correlation

Adapted from Torralba



# Classification: Simple Correlation



Simple template matching is not going to make it

Adapted from Torralba

# Classification Challenges



# Simple Classification

## ■ Nearest neighbor

No training required!

Training  
examples  
from class 1

Test  
example



Training  
examples  
from  
class 2

### ■ Pros:

- + Simple to implement
- + Can be any distribution
- + Works for any number of classes
- + *Nonparametric* method

### ■ NN cons:

- Need good distance function
- Slow at test time





Less  
similar



More  
similar



# Simple Correlation

- Expected to fail to generalize to new
  - Views
  - Illumination
  - Non rigid transformation
  - Occlusion
  - ...

# Main Challenges

- Generalization:  
Recognition in unseen images
- Robustness
- Scalability:
  - Space and time



# Object Recognition Evolution

- Tailored end-to-end
  - Features (e.g., SIFTS, HOGS, RGB histograms,...)
  - Mapping (e.g., nearest neighbor, majority, ...)
- Learning Half way
  - Features (e.g., BOW)
  - Mapping (e.g., SVM)
- Learning end-to-end
  - Features & mapping (e.g., neural networks)

# Examples

## Tailored End-to-End

- Simple image correlation using nearest neighbor or k-nearest neighbors
- Tailored facial feature representation (Yuille 1991)
- General image features, e.g., SIFTS (Lowe 2004)





Figure 3. The eye template.



# Machine Learning

- Use a set of observation to uncover the relation between input and the desired output
  - The alternative: study the problem mathematically
- Types:
  - Supervised/Semi-supervised/unsupervised

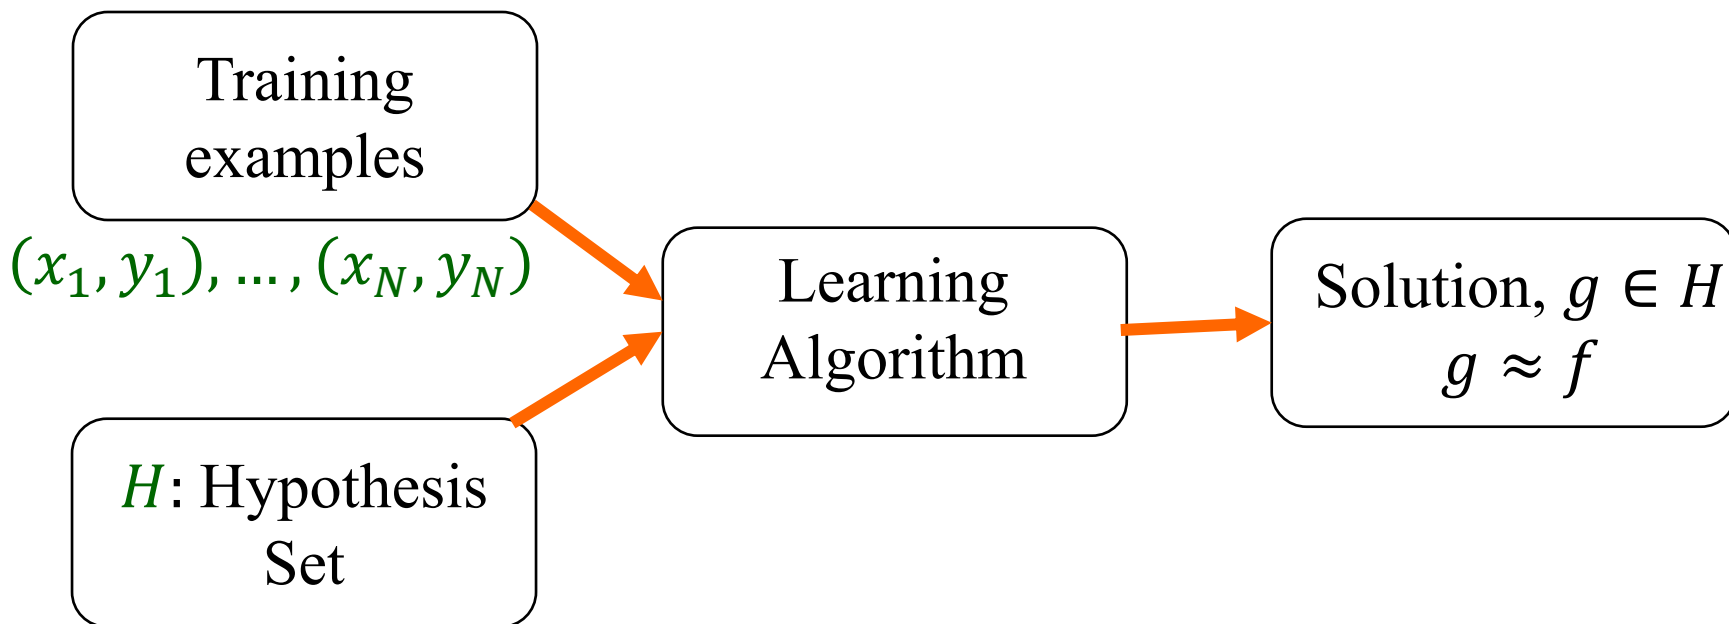
# The Essences of Learning

- A pattern exists:
  - E.g., a picture of the same object has something in common
- Close mathematical solution is unknown
  - Otherwise – use it!
- Must have data
  - Labeled
  - Unlabeled

# Components of General Learning

- Input:  $x \in X$
- Output:  $y \in Y$
- Target function:  $f: X \rightarrow Y$
- Training examples:  $(x_1, y_1), \dots, (x_N, y_N)$
- Hypothesis:  $g: X \rightarrow Y$ , s.t.,  $g \approx f$

# Outline for Learning



- Target function:  $f: X \rightarrow Y$
- Find  $g: X \rightarrow Y$ , s.t.,  $g \approx f$

# Object Recognition

- The function  $f$ : determine the object in the image
- Learn  $f$  from a set of labeled/unlabeled examples
- Use  $f$  to recognize objects in a new image

# Supervised Learning


Training:



Positive examples




Negative examples






# Supervised Learning

Testing: face?



# Supervised Learning

Testing: face?



Black Box




no

# Learning Half Way

- Separate representation & matching
- Main Questions:
  - What is the representation
  - What is the “black box”

# Tailored Representation

- The entire image
- Histogram
- Contours
- Features:
  - Regular grid
  - Local interest points
- Descriptors
  - Patches, Histograms, SIFT...




# “Black Box”

- Examples:
  - Fisher classifier
  - SVM
  - Random Forest
  - ...
  - Neural networks
  - Convolution neural network (CNN)

# Bag of Words

## (Sivic & Zisserman 2003)

- Use matching techniques from document analysis



**Object**

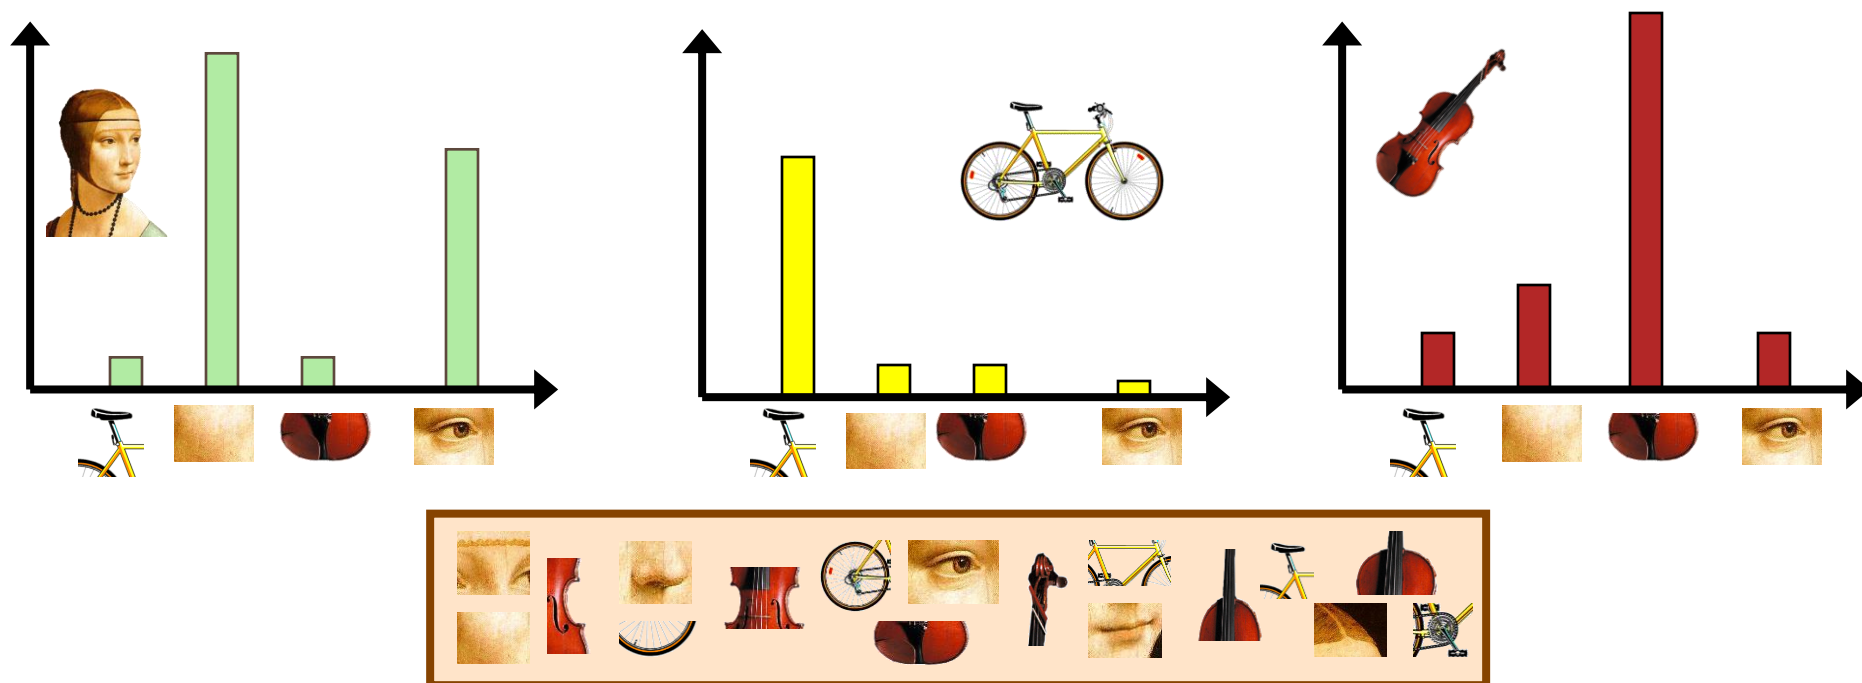


**Bag of 'words'**





# Bag of Words

- Independent features
- Representation: histogram of features






# Collect Features



## 2. Codewords dictionary formation




# K-Means

- Choose a fixed number of clusters
- Find the centers of the clusters
- Minimize: 
$$\sum_{i \in \text{clusters}} \left\{ \sum_{j \in \text{elements of } i\text{'th cluster}} \|x_j - \mu_i\|^2 \right\}$$

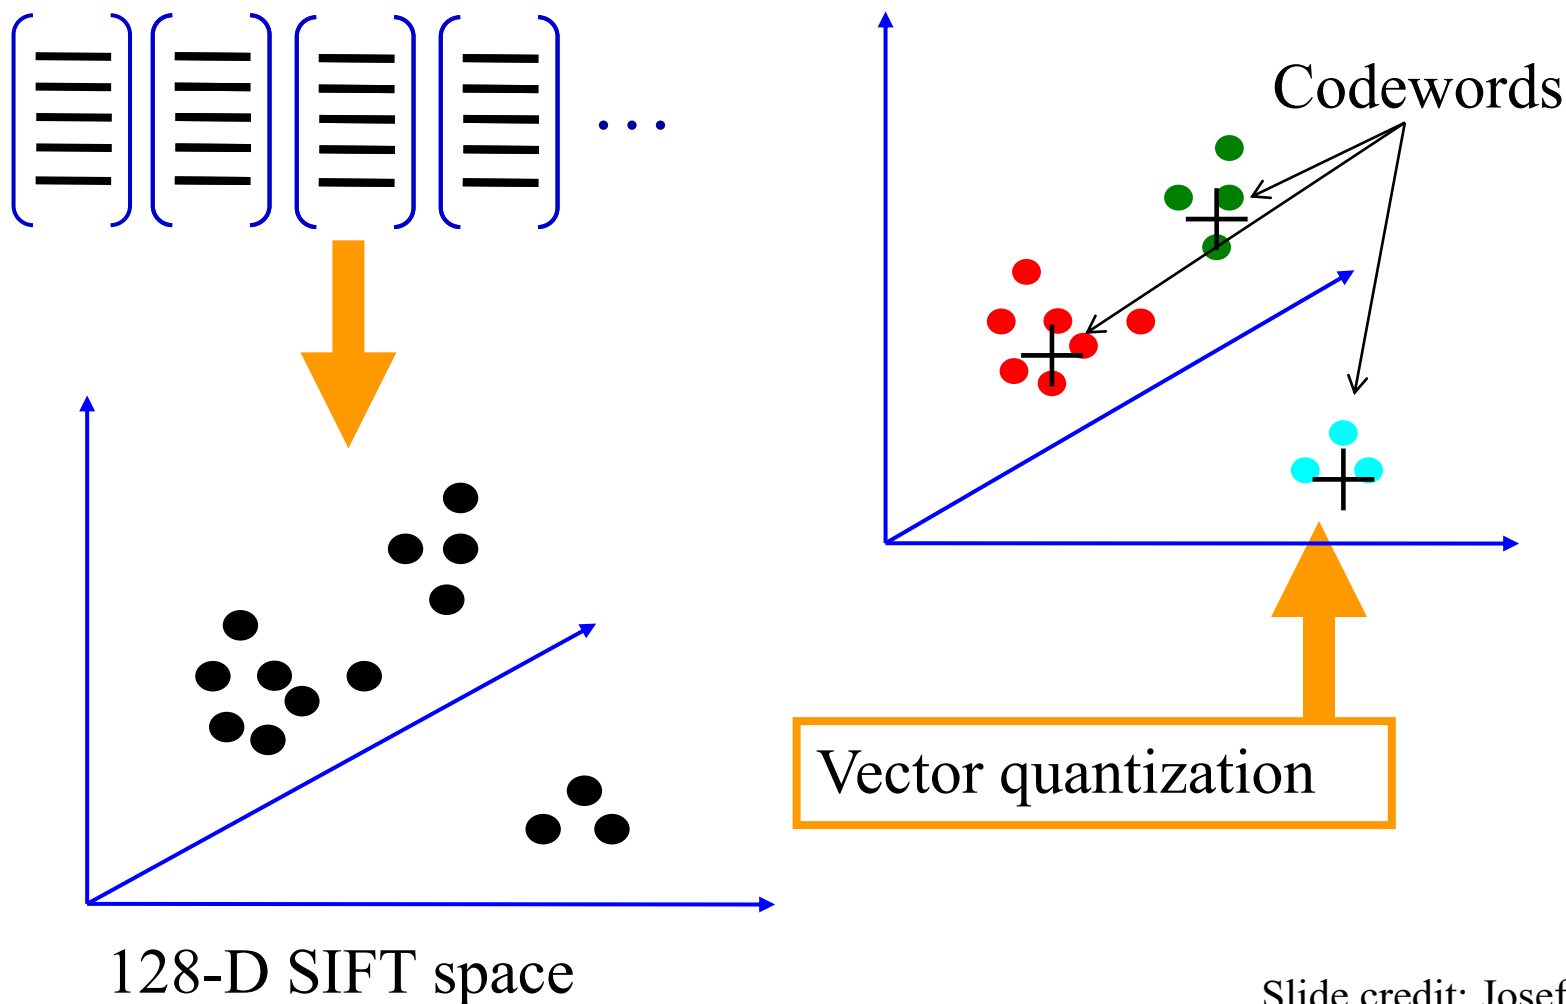
where  $\mu_i$  is the center of  $S_i$

- Iterative Algorithm
  - Fix cluster centers; allocate points to closest cluster
  - Fix allocation; compute best cluster centers:

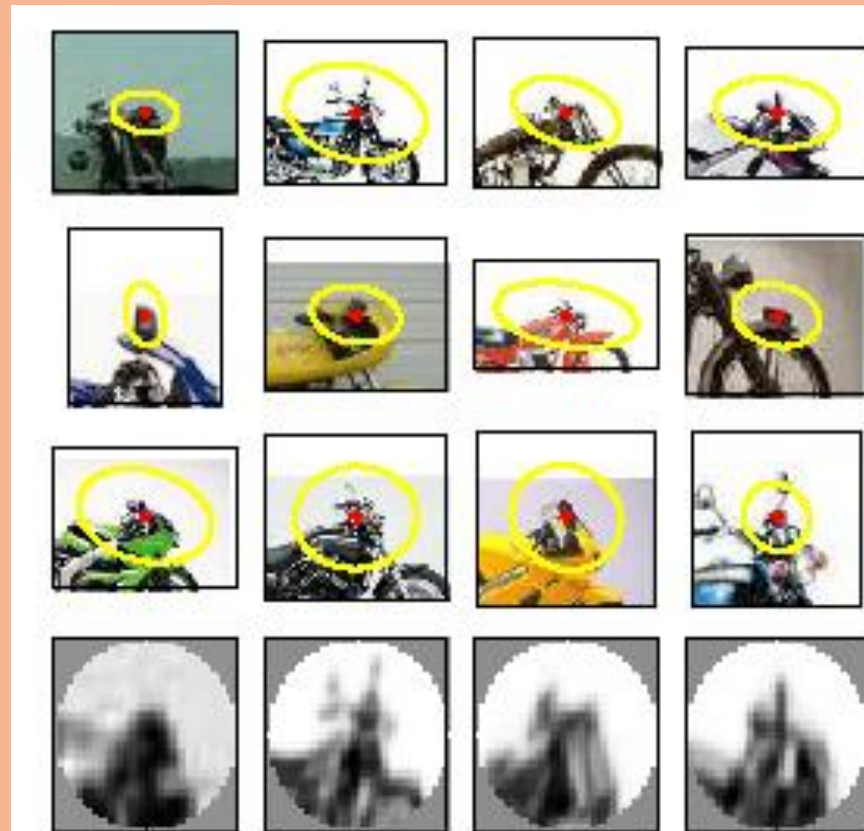

# K-mean Clustering



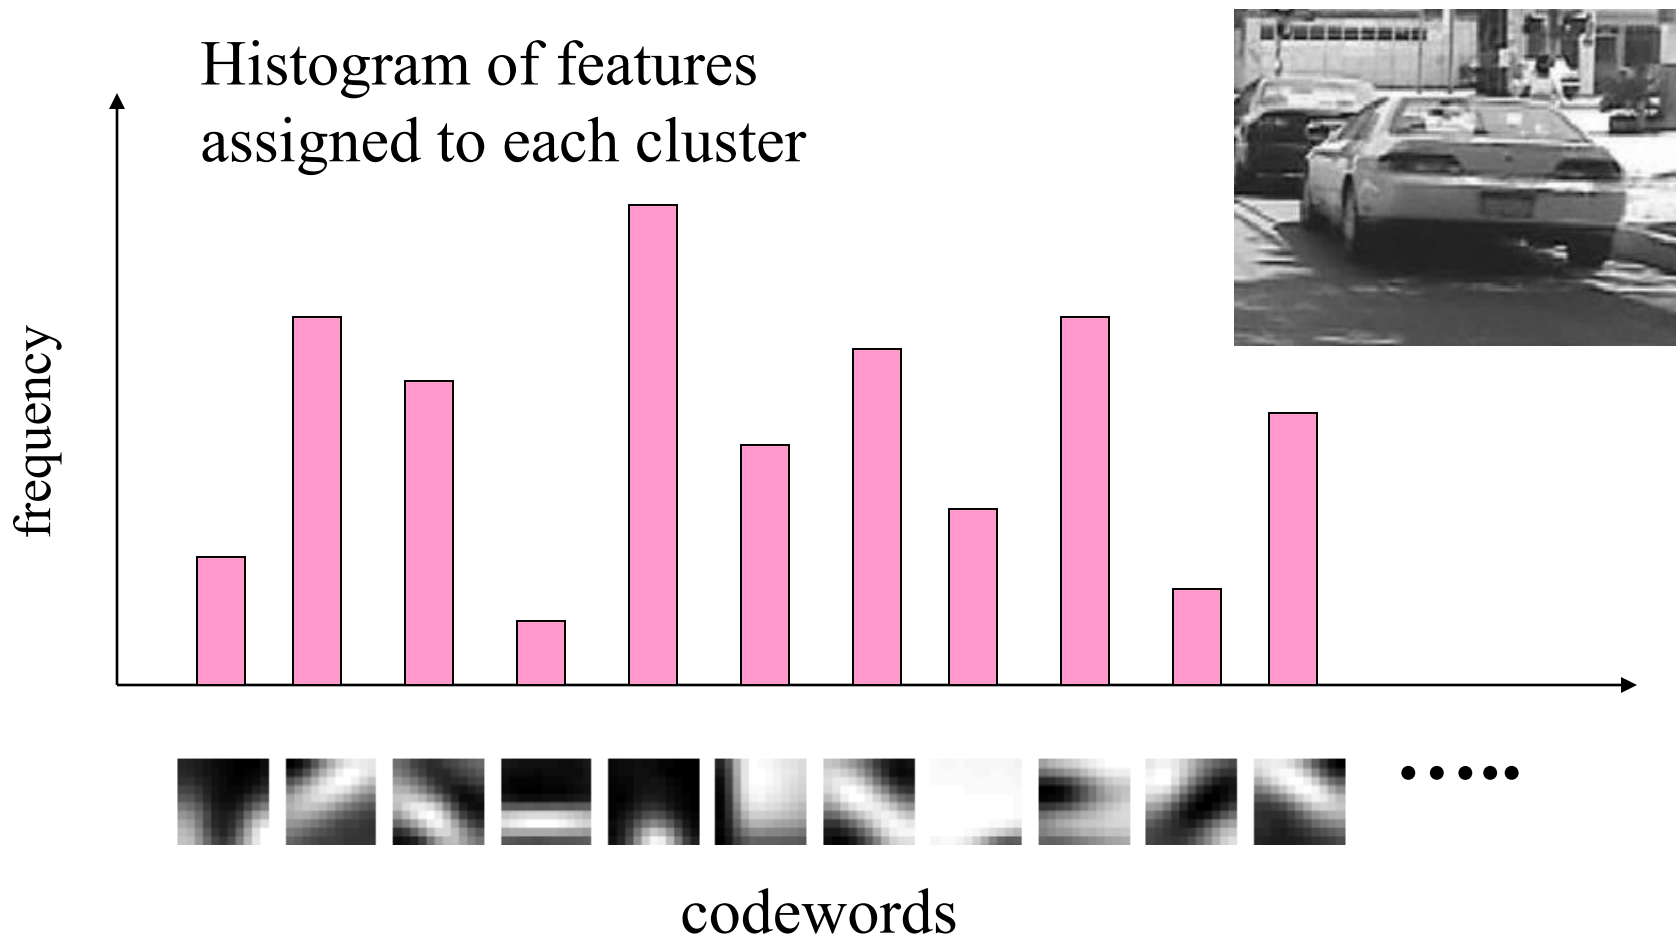
## 2. Codewords dictionary formation



# Image patch examples of codewords



# Image Representation





# Weighted Histogram

- Weight a word according to its frequency in all documents




# Uses of BoW Representation

- Treat as feature vector for standard classifier
  - e.g., SVM
- Cluster BoW vectors over image collection 
  - Discover visual themes
- Hierarchical models 
  - Decompose scene/object

# Support Vector Machine (SVM)



- Linear classifier



# Limitations?



# Other Classic Recognition Methods

- Geometry:
  - Invariance to view points
  - 3D alignment
  - Weak geometric constraints

# Summary

- Goals of computer vision:
  - Infer the physical world from images
  - Impart human perceptual ability to machines
  - Improve over human perceptual ability

# What we learned

- What are the challenges
- Tasks examples
- Example of solutions (algorithms)
- Principals of solutions
- Implementations: parallelism

# Course Topics

- Image features:
  - edges, corners, other interest points
- Image formation: Geometry
- Geometry: Stereo, SFM, Homography, epipolar geometry
- Motion analysis:
  - Optical flow, change detection, tracking
- Object Recognition
  - BOW

# Tools

- Convolution
- Algebra:
  - Projective Algebra, SVD, 2D/3D transformations, Use of EigenVectors,
- Gaussians and Mixture of Gaussians
- RANSAC
- Evaluation methods



# More Tools

- Double threshold
- RGB Histograms & histogram of gradients
- Nearest neighbor & Lowe distance
- Use vectorized solutions!
- ...


# Left for Additional Lectures

- More on object recognition
  - Supervised/unsupervised, AdaBoost, Cascade
- More on multi camera
- Segmentation
- Photometry:
  - Image formation
  - Photometric Stereo
- Introduction to CNN
  - state-of-the-art compared to classic methods

# Depth Perception




# Depth Perception





# Shape from Shading is ill-pose





# CrowdCam







# Papal coronation






# CrowdCam: Moving Object Segmentation





# Depth Ambiguity

- Illusion





# Thanks

See you in the one-on-one meeting