

Statistics and data analysis

Zohar Yakhini, Leon Anavy, Ben Galili


IDC, Herzeliya

Distributions and a brief review of basic concepts

0001010110111100101010100101001110001010111110010101010010100
1011010010100100000111100101010010001010010101001111101100101101001
01010010100101010010100001010111001010101001010011100010100010
1011111001010101001010010101001010010100100100000111100101001000100010
10100111110110010110010101001010010101001010100100100100100101011
0101010100101111010011110010101010101010101001010101010011110010
1110010101001000101001010100111101100101100101010101001010101010
010101001001100101010101010101010101010111001010101010101001010
0100001111001010100100010100101010111110100101100101010010101010
0101010010101010010011001010101010111111010011110010101011111
1110110010110010101001010010101010101010100001010111001010101001
010011010011100010101111100101010101010101001010001000001111100
101010010001010010101010101111101001011100101001010101001010
10100100110010101101010101010101010100111001010101010111010011010
10101010011110010000101010100110100111001010101010100010001111110
11100101010010001010010101010111101100101101001011100101010010100
1101001110001010111110010101010010100101001010010000011111001010
1001000101001010101010111110110010110110010101001010101010101010
0100110010
1010011111001011001010100100111110101010101010101010101010101001
0101001010101001010101010101010101010101111001011100101010101010
11010101001010010000011111010011111010010010001111101001111101100



0010011101010100101010100100100010
1010100010101111101011010011001001
1110101010011010010110010110110010



Basics: (Ω, P)

Ω = the probability sample space

A collection (an algebra) of measurable events – sets of samples


A probability measure, P – assigns a probability to every measurable set of samples

The measure P is additive for disjoint sets, it's non-negative and $P(\Omega) = 1$

Example – Rolling 2 Dice (Red/Green)

Ω = All possible outcomes

Measurable sets = all subsets of this finite space



Red\Green	1	2	3	4	5	6
1	1,1	1,2	1,3	1,4	1,5	1,6
2	2,1	2,2	2,3	2,4	2,5	2,6
3	3,1	3,2	3,3	3,4	3,5	3,6
4	4,1	4,2	4,3	4,4	4,5	4,6
5	5,1	5,2	5,3	5,4	5,5	5,6
6	6,1	6,2	6,3	6,4	6,5	6,6

Random Variables

- A Random Variable (RV) is a numerical function defined on the probability sample space.
- For each element of a sample space, the random variable takes on exactly one value
- Random Variables are usually denoted using upper case letters (X, Y)
- Individual outcomes for an RV are usually denoted using lower case letters (x, y)

Example:

- Toss a coin 5 times.
- The sample space –
 Ω = all possible outcomes:
00000, 00001, 00010, 00011, ... , 01111,
..., 11101, 11110, 11111
- One possible RV defined on this space – the outcome of the third toss
- Another : Y = total number of 1s
(what is $P(Y = 1)$ in this case?)
- Is the number of 1s prime? (a binary RV)
- Another : count how many 1s on even numbered tosses
- Count how many 11 runs

Example – Rolling 2 Dice (Red/Green)

Ω = All possible outcomes



Red\Green	1	2	3	4	5	6
1	2	3	4	5	6	7
2	3	4	5	6	7	8
3	4	5	6	7	8	9
4	5	6	7	8	9	10
5	6	7	8	9	10	11
6	7	8	9	10	11	12

Probability Distributions

- Probability Distribution: Table, Graph, or Formula that describes values a random variable can take on, and their corresponding occurrence probabilities (discrete RV) or density (continuous RV)
- Discrete Probability Distribution: Assigns probabilities (masses) to the individual outcomes
- Continuous Probability Distribution: Assigns density at individual points, probability of ranges can be obtained by integrating density function
- Discrete probability distribution: $p(y) = P(Y = y)$
- Continuous densities are denoted by $f(y)$.
We then have

$$P(Y \in I) = \int_I f(y) dy$$

- Cumulative Distribution Function: $F(y) = P(Y \leq y)$
- Probability distributions sum or integrate to 1.
- What values can the CDF, F , of a random variable take?


Sum of 2 Dice – Probability Mass Function & CDF

y	p(y)	F(y)
2	1/36	1/36
3	2/36	3/36
4	3/36	6/36
5	4/36	10/36
6	5/36	15/36
7	6/36	21/36
8	5/36	26/36
9	4/36	30/36
10	3/36	33/36
11	2/36	35/36
12	1/36	36/36

$$p(y) = \frac{\text{\# of ways 2 dice can sum to } y}{\text{\# of ways 2 dice configurations}}$$

$$F(y) = \sum_{t=2}^y p(t)$$


Rolling 2 Dice – Probability Mass Function



Expected Values of Discrete RV's

- Mean (aka Expected Value) – the weighted average value an RV (or function of RV). Weighting is according to the underlying probability space.
- Variance – Average squared deviation between a realization of an RV (or function of RV) and its mean
- Standard Deviation – Positive Square Root of Variance (in same units as the data)
- Notation:
 - Mean: $E(Y) = \mu$
 - Variance: $\text{Var}(Y) = \sigma^2$
 - Standard Deviation: σ

$$E(X) = \sum_{\text{all relevant } x} x p(x)$$



How much will we pay (or not) to play this game?

Expected Value and Variance of Discrete RV's

Mean : $E(Y) = \mu = \sum_{\text{all } y} yp(y)$

Mean of a function $g(Y)$: $E[g(Y)] = \sum_{\text{all } y} g(y)p(y)$

Variance : $V(Y) = \sigma^2 = E[(Y - E(Y))^2] = E[(Y - \mu)^2] =$
 $= \sum_{\text{all } y} (y - \mu)^2 p(y) = \sum_{\text{all } y} (y^2 - 2y\mu + \mu^2)p(y) =$
 $= \sum_{\text{all } y} y^2 p(y) - 2\mu \sum_{\text{all } y} yp(y) + \mu^2 \sum_{\text{all } y} p(y) =$
 $= E[Y^2] - 2\mu(\mu) + \mu^2(1) = E[Y^2] - \mu^2$

Standard Deviation : $\sigma = +\sqrt{\sigma^2}$

Expected Values of Linear Functions of Discrete RV's

Linear Functions : $g(Y) = aY + b$ ($a, b \equiv \text{constants}$)

$$E[aY + b] = \sum_{\text{all } y} (ay + b)p(y) =$$

$$= a \sum_{\text{all } y} yp(y) + b \sum_{\text{all } y} p(y) = a\mu + b$$

$$V[aY + b] = \sum_{\text{all } y} ((ay + b) - (a\mu + b))^2 p(y) =$$

$$\sum_{\text{all } y} (ay - a\mu)^2 p(y) = \sum_{\text{all } y} [a^2 (y - \mu)^2] p(y) =$$

$$= a^2 \sum_{\text{all } y} (y - \mu)^2 p(y) = a^2 \sigma^2$$

$$\sigma_{aY+b} = |a|\sigma$$

Example – Rolling 2 Dice

y	p(y)	yp(y)	y ² p(y)
2	1/36	2/36	4/36
3	2/36	6/36	18/36
4	3/36	12/36	48/36
5	4/36	20/36	100/36
6	5/36	30/36	180/36
7	6/36	42/36	294/36
8	5/36	40/36	320/36
9	4/36	36/36	324/36
10	3/36	30/36	300/36
11	2/36	22/36	242/36
12	1/36	12/36	144/36
Sum	36/36 =1.00	252/36 =7.00	1974/36= 54.833



$$\mu = E(Y) = \sum_{y=2}^{12} yp(y) = 7.0$$

$$\begin{aligned}\sigma^2 &= E[Y^2] - \mu^2 = \sum_{y=2}^{12} y^2 p(y) - \mu^2 \\ &= 54.8333 - (7.0)^2 = 5.8333\end{aligned}$$

$$\sigma = \sqrt{5.8333} = 2.4152$$

Expectation - another angle

Consider a probability space (Ω, P) and a rv $X: \Omega \rightarrow \mathbb{R}$

An equivalent definition of the expected value is:

$$E(X) = \sum_{\omega \in \Omega} X(\omega)P(\omega)$$

A very important conclusion is:

$$\begin{aligned} E(X + Y) &= \sum_{\omega \in \Omega} (X(\omega) + Y(\omega))P(\omega) \\ &= \sum_{\omega \in \Omega} X(\omega)P(\omega) + \sum_{\omega \in \Omega} Y(\omega)P(\omega) \\ &= E(X) + E(Y) \end{aligned}$$



Linearity of expectations

Red\Green	1	2	3	4	5	6
1	2	3	4	5	6	7
2	3	4	5	6	7	8
3	4	5	6	7	8	9
4	5	6	7	8	9	10
5	6	7	8	9	10	11
6	7	8	9	10	11	12

$$\forall \lambda > 0 \quad P(|X - \mu| > \lambda) \leq \frac{V(X)}{\lambda^2}$$

Deviation from the mean

- Tchebysheff's theorem: Suppose Y is any random variable with mean μ and standard deviation σ . Then:

$$P(\mu - b\sigma \leq Y \leq \mu + b\sigma) \geq 1 - (1/b^2) \text{ for } b > 0$$

- $b=1$: $P(\mu - 1\sigma \leq Y \leq \mu + 1\sigma) \geq 1 - (1/1^2) = 0$ (trivial result)
 - $b=2$: $P(\mu - 2\sigma \leq Y \leq \mu + 2\sigma) \geq 1 - (1/2^2) = 3/4$
 - $b=3$: $P(\mu - 3\sigma \leq Y \leq \mu + 3\sigma) \geq 1 - (1/3^2) = 8/9$
- Note that this is a very conservative bound, but that it works for any distribution

- For Mound Shaped Distributions, aka Gaussians:

- $k=1$: $P(\mu - 1\sigma \leq Y \leq \mu + 1\sigma) \approx 0.68$
- $k=2$: $P(\mu - 2\sigma \leq Y \leq \mu + 2\sigma) \approx 0.95$
- $k=3$: $P(\mu - 3\sigma \leq Y \leq \mu + 3\sigma) \approx 0.995$

Proof of Tchebysheff's Theorem

Breaking real line into 3 parts :

i) $(-\infty, (\mu - k\sigma)^-]$ ii) $[(\mu - k\sigma), (\mu + k\sigma)]$ iii) $[(\mu + k\sigma)^+, \infty)$

Making use of the definition of Variance :

$$V(Y) = \sigma^2 = \sum_{-\infty}^{\infty} (y - \mu)^2 p(y) =$$

$$\sum_{-\infty}^{(\mu - k\sigma)^-} (y - \mu)^2 p(y) + \sum_{(\mu - k\sigma)}^{(\mu + k\sigma)} (y - \mu)^2 p(y) + \sum_{(\mu + k\sigma)^+}^{\infty} (y - \mu)^2 p(y)$$

In Region i): $y - \mu \leq -k\sigma \Rightarrow (y - \mu)^2 \geq k^2 \sigma^2$

In Region iii): $y - \mu \geq k\sigma \Rightarrow (y - \mu)^2 \geq k^2 \sigma^2$

$$\Rightarrow \sigma^2 \geq k^2 \sigma^2 P(Y < \mu - k\sigma) + \sum_{(\mu - k\sigma)}^{(\mu + k\sigma)} (y - \mu)^2 p(y) + k^2 \sigma^2 P(Y > \mu + k\sigma)$$

$$\begin{aligned} \Rightarrow \sigma^2 &\geq k^2 \sigma^2 P(Y < \mu - k\sigma) + k^2 \sigma^2 P(Y > \mu + k\sigma) = \\ &= k^2 \sigma^2 [1 - P(\mu - k\sigma \leq Y \leq \mu + k\sigma)] \end{aligned}$$

$$\Rightarrow \frac{\sigma^2}{k^2 \sigma^2} = \frac{1}{k^2} \geq [1 - P(\mu - k\sigma \leq Y \leq \mu + k\sigma)] \Rightarrow P(\mu - k\sigma \leq Y \leq \mu + k\sigma) \geq 1 - \frac{1}{k^2}$$

Discrete Uniform Distribution

- Suppose Y can take on any integer value between a and b inclusive, each equally likely (e.g. rolling a dice, where $a=1$ and $b=6$). Then Y follows the discrete uniform distribution.

$$f(y) = \frac{1}{b - (a - 1)} \quad a \leq y \leq b$$

$$F(y) = \begin{cases} 0 & y < a \\ \frac{\text{int}(y) - (a - 1)}{b - (a - 1)} & a \leq y < b \\ 1 & y \geq b \end{cases} \quad \text{int}(x) \equiv \text{integer portion of } x$$

$$E(Y) = \sum_{y=a}^b y \left(\frac{1}{b - (a - 1)} \right) = \frac{1}{b - (a - 1)} \left[\sum_{y=1}^b y - \sum_{y=1}^{a-1} y \right] = \frac{1}{b - (a - 1)} \left[\frac{b(b+1)}{2} - \frac{(a-1)a}{2} \right] = \frac{b(b+1) - a(a-1)}{2(b - (a - 1))}$$

$$\begin{aligned} E(Y^2) &= \sum_{y=a}^b y^2 \left(\frac{1}{b - (a - 1)} \right) = \frac{1}{b - (a - 1)} \left[\sum_{y=1}^b y^2 - \sum_{y=1}^{a-1} y^2 \right] = \frac{1}{b - (a - 1)} \left[\frac{b(b+1)(2b+1)}{6} - \frac{(a-1)a(2a-1)}{6} \right] \\ &= \frac{b(b+1)(2b+1) - a(a-1)(2a-1)}{6(b - (a - 1))} \end{aligned}$$

$$\Rightarrow V(Y) = E(Y^2) - [E(Y)]^2 = \frac{b(b+1)(2b+1) - a(a-1)(2a-1)}{6(b - (a - 1))} - \left[\frac{b(b+1) - a(a-1)}{2(b - (a - 1))} \right]^2$$

Note : When $a = 1$ and $b = n$:

$$E(Y) = \frac{n+1}{2} \quad V(Y) = \frac{(n+1)(n-1)}{12} \quad \sigma = \sqrt{\frac{(n+1)(n-1)}{12}}$$

Bernoulli Distribution

- An experiment consists of one trial. It can result in one of 2 outcomes: Success or Failure (or a property being Present or Absent).
- Probability of Success ($Y = 1$) is p ($0 < p < 1$)
- Example: coin tossing


$$p(y) = \begin{cases} p & y = 1 \\ 1 - p & y = 0 \end{cases}$$

$$E(Y) = \sum_{y=0}^1 yp(y) = 0(1 - p) + 1p = p$$


$$E(Y^2) = 0^2(1 - p) + 1^2 p = p$$

$$\Rightarrow V(Y) = E(Y^2) - [E(Y)]^2 = p - p^2 = p(1 - p)$$

$$\Rightarrow \sigma = \sqrt{p(1 - p)}$$



Binomial Distribution



- A binomial experiment consists of a series of n identical trials
- Consider all possible tossing trajectories. This is our probability space, Ω .
- Each trial is Bernoulli as above
- Trials are independent (outcome of one has no bearing on outcomes of others – formal definition next week)
- Probability of Success, p , is constant for all trials
- The random variable Y which counts the number of Successes in the n trials is said to follow a **Binomial Distribution with parameters n and p**
- Y can take on the values $y = 0, 1, \dots, n$
- Notation: $Y \sim \text{Binom}(n, p)$


Binomial Distribution

- $P(Y = k) = \binom{n}{k} p^k (1 - p)^{n-k}$
- $\sum_{k=0}^n P(Y = k) = \sum_{k=0}^n \binom{n}{k} p^k (1 - p)^{n-k} = (p + 1 - p)^n = 1$


$\omega \in \Omega :$




Binom(5,0.5) –
tossing a fair coin 5 times, counting successes




Binom(5,0.3)




Binom(10,0.5) –
tossing a fair coin 10 times, counting successes




Binom(50,0.5) –
tossing a fair coin 50 times, counting successes



Tossing 10 coins w $p = 0.1, 0.5, 0.8$



Tossing 50 coins w $p = 0.1, 0.5, 0.8$



Binomial Distribution – Expected Value

$$f(y) = \frac{n!}{y!(n-y)!} p^y q^{n-y} \quad y = 0, 1, \dots, n \quad q = 1 - p$$

$$E(Y) = \sum_{y=0}^n y \left[\frac{n!}{y!(n-y)!} p^y q^{n-y} \right] = \sum_{y=1}^n y \left[\frac{n!}{y!(n-y)!} p^y q^{n-y} \right]$$

(Summand = 0 when $y = 0$)

$$\Rightarrow E(Y) = \sum_{y=1}^n \left[\frac{yn!}{y(y-1)!(n-y)!} p^y q^{n-y} \right] = \sum_{y=1}^n \left[\frac{n!}{(y-1)!(n-y)!} p^y q^{n-y} \right]$$

Let $y^* = y - 1 \Rightarrow y = y^* + 1$ Note: $y = 1, \dots, n \Rightarrow y^* = 0, \dots, n - 1$

$$\begin{aligned} \Rightarrow E(Y) &= \sum_{y^*=0}^{n-1} \frac{n(n-1)!}{y^*!(n-(y^*+1))!} p^{y^*+1} q^{n-(y^*+1)} = np \sum_{y^*=0}^{n-1} \frac{(n-1)!}{y^*!((n-1)-y^*)!} p^{y^*} q^{(n-1)-y^*} = \\ &= np(p+q)^{n-1} = np(p+(1-p))^{n-1} = np(1) = np \end{aligned}$$

$\omega \in \Omega :$



Binomial Distribution – Variance and S.D.

$$f(y) = \frac{n!}{y!(n-y)!} p^y q^{n-y} \quad y = 0, 1, \dots, n \quad q = 1 - p$$

Note: $E(Y^2)$ is difficult (impossible?) to get, but $E(Y(Y-1)) = E(Y^2) - E(Y)$ is not:

$$E(Y(Y-1)) = \sum_{y=0}^n y(y-1) \left[\frac{n!}{y!(n-y)!} p^y q^{n-y} \right] = \sum_{y=2}^n y(y-1) \left[\frac{n!}{y!(n-y)!} p^y q^{n-y} \right]$$

(Summand = 0 when $y = 0, 1$)

$$\Rightarrow E(Y(Y-1)) = \sum_{y=2}^n \frac{n!}{(y-2)!(n-y)!} p^y q^{n-y}$$

Let $y^{**} = y - 2 \Rightarrow y = y^{**} + 2$ Note: $y = 2, \dots, n \Rightarrow y^{**} = 0, \dots, n-2$

$$\Rightarrow E(Y(Y-1)) = \sum_{y^{**}=0}^{n-2} \frac{n(n-1)(n-2)!}{y^{**}!(n-(y^{**}+2))!} p^{y^{**}+2} q^{n-(y^{**}+2)} = n(n-1)p^2 \sum_{y^{**}=0}^{n-2} \frac{(n-2)!}{y^{**}!((n-2)-y^{**})!} p^{y^{**}} q^{(n-2)-y^{**}} =$$

$$= n(n-1)p^2 (p+q)^{n-2} = n(n-1)p^2 (p+(1-p))^{n-2} = n(n-1)p^2$$

$$\Rightarrow E(Y^2) = E(Y(Y-1)) + E(Y) = n(n-1)p^2 + np = np[(n-1)p + 1] = n^2 p^2 - np^2 + np = n^2 p^2 + np(1-p)$$

$$\Rightarrow V(Y) = E(Y^2) - [E(Y)]^2 = n^2 p^2 + np(1-p) - (np)^2 = np(1-p)$$

$$\Rightarrow \sigma = \sqrt{np(1-p)}$$

Using the binomial distribution.

Example: Experimental treatment for Kidney Cancer

- Suppose we have $n = 40$ patients who will be receiving an experimental therapy (Tx) which is believed to be better than current treatments (standard of care = SoC). The latter has a historically derived 5-year survival rate of 20%. That is, under the SoC the probability of 5-year survival is $p = 0.2$
- We will now count 5-year survival under Tx and will then need to decide if we can confidently say that the new experimental treatment is better.

Results and “The Question”

- Suppose that using the new treatment we find that 16 out of the 40 patients survive at least 5 years past diagnosis.
- Q: Does this result suggest that the new therapy, Tx, has a better 5-year survival rate than that of the SoC?
That is:
is the probability that a patient survives at least 5 years greater than 0.2 when treated using the new therapy?

What do we consider in answering the question of interest?

We essentially ask ourselves the following:

- If we assume that new therapy is **no better** than the current then what is the probability of seeing the observed numbers? That is – how likely are they to occur, in such case, by chance alone?
- More specifically:
What is the probability of seeing 16 or more successes out of 40 if the success rate of the new therapy is also 0.2?
- This is called estimating the **p-value** of the **OBSERVED RESULT** under the **NULL model**

Binomial

- This is a binomial experiment situation...
 - There are $n = 40$ patients and we are counting the number of patients that survive 5 or more years.
 - The individual patient outcomes are independent and under the NULL MODEL the probability of success is $p = 0.2$ for all patients.
(that is: we assume that Tx is NOT better than the standard of care)
- So the random variable $X = \# \text{ of "successes" in the clinical trial}$ is, under the NULL model, Binomial with $n = 40$ and $p = 0.2$,

i.e., under the null: $X \sim \text{Binomial}(40, 0.2)$

Example: Treatment of Kidney Cancer - cont

- $X \sim \text{BIN}(40, 0.2)$, find the probability that exactly 16 patients survive at least 5 years.

$$P(X = 16) = \binom{40}{16} \cdot 20^{16} \cdot 80^{24} = .001945$$

- This requires some calculator gymnastics and some scratchwork (or a Matlab command ...)
- But - keep in mind that we need to find the probability of having **16 or more** patients surviving at least 5 yrs.

Example: Treatment of Kidney Cancer

- So we actually need to find:

$$P(X \geq 16) = P(X = 16) + P(X = 17) + \dots + P(X = 40)$$

$$P(X = 16) = \binom{40}{16} \cdot 20^{16} \cdot 80^{24} = .001945$$

+

$$P(X = 17) = \binom{40}{17} \cdot 20^{17} \cdot 80^{23} = .000686$$

...

+

$$P(X = 40) = \binom{40}{40} \cdot 20^{40} \cdot 80^0 \approx 0$$

$$= .002936 \quad \text{Yupp!}$$

When using commands in a statistical language we will use the CDF

```
from scipy.stats import binom
rv = binom(40, 0.2)
x_16_and_up = 1 - rv.cdf(15)
print("{:.4f}".format(x_16_and_up))

0.0029
```

Conclusion (statistics helps decision making ...)

Because it is highly unlikely ($p = 0.0029$) that we would see this many successes in a group of 40 patients if the new Tx had the same probability of success as the SoC we have to make a choice, either ...

A) Tx's survival rate is less than 0.2 and we have obtained a very rare result by chance.

OR


B) our assumption about the success rate of the new Tx is wrong and in actuality it has a better than 20% 5-year survival rate making the observed result more plausible.

Caveat: other aspects of the null model can also be wrong ...

Tx is better than the
SoC treatment with
p-value <0.003 under
a binomial null model

Next week we will start with two waiting time distributions: the geometric distribution and the negative binomial (Polya) distribution

$\omega \in \Omega :$



Continue to infinity ...

A Geometric random variable:

$X(\omega) = \text{time of first success}$

Summary and what's next

- Statistics provides tools and frameworks for the rigorous interpretation of data, for effective (and efficient) inference and for clearly presenting and stating conclusions.
- Data analysis uses computational approaches to implement statistical principles in practically analyzing data.
- In this course we will address theoretical and practical aspects of both.
- We will emphasize computer age aspects: efficiency, volume etc
- We learned about Bernoulli random variables and about the Binomial distribution.
- Next time: Geometric, NegB, Poisson distributions and related aspects
- During the course we will present and investigate more distributions.
- We proved Tchebychef's Thm and saw how it yields a bound on large deviations from the mean
- In following weeks we will derive more efficient approaches/bounds and see how to use them in practice.
- Next week: independence or not? And the consequences ...