

Markov Chains


Statistics and data analysis

Ben Galili



Zohar Yakhini

Leon Anavy

IDC, Herzeliya



0010011101010100101010100100100010
1010100010101111101011010011001001
1110101010011010010110010110110010



Stochastic Processes

A (discrete, integer indexed) stochastic process is a sequence of random variables $X_0, X_1, X_2, \dots, X_n, \dots$ with a joint distribution defined over any finite set of variables.

For example, the variables: $(X_8, X_{11}, X_{23}, X_{30}, X_{31}, X_{207})$ have some 6-dimensional probability distribution which is consistent with the (lower-dimensional) probability distributions of all subsets of random variables therein.

Example: independent coin tossing.

Stationary Stochastic Processes

A process is called stationary if for every t the distribution of any finite set of variables $(X_{i1}, X_{i2}, X_{i3}, X_{i4} \dots, X_{ir})$ is the same as the distribution of $(X_{i1+t}, X_{i2+t}, X_{i3+t}, X_{i4+t} \dots, X_{ir+t})$

A process is temporally homogeneous if for every t the distribution of any finite set of variables $(X_{i1}, X_{i2}, X_{i3}, X_{i4} \dots, X_{ir} | X_{i0})$ is the same as the distribution of $(X_{i1+t}, X_{i2+t}, X_{i3+t}, X_{i4+t} \dots, X_{ir+t} | X_{i0+t})$

Is independent coin tossing w $p = 0.5$ stationary? How about with $p = 0.3$?

Markov chains

$X_0, X_1, X_2, \dots, X_n, \dots$
s.t. $X_i \in \{\text{Sunny}, \text{Cloudy}, \text{Rainy}\}$

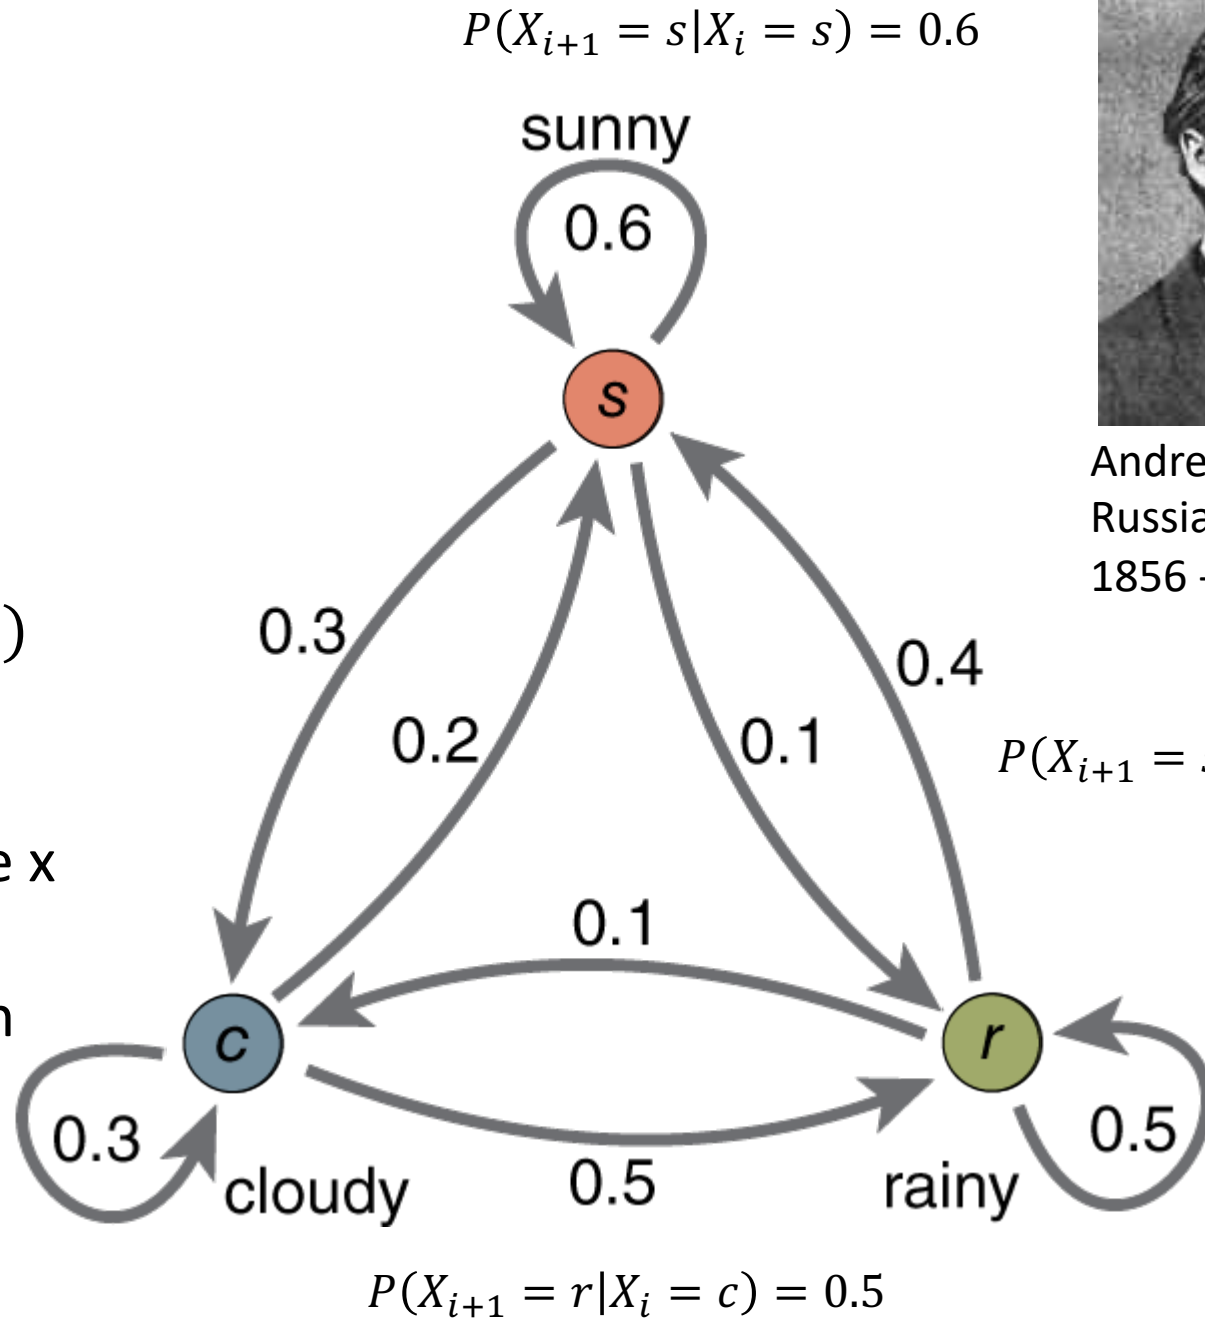
Example sets of outcomes:

$(X_8 = s, X_{11} = s, X_{23} = r, X_{30} = r)$

Transition probability:

The probability of moving from state x to state y

The probability distribution of step n
only depend on step $n-1$



Andrey A Markov
Russian Mathematician
1856 - 1922

Markov chains

A stochastic process $X_0, X_1, X_2, \dots, X_n, \dots$

Taking values in a state space $\{S_1, \dots, S_k, \dots\}$

Is said to be a (finite state space) Markov process if:

- It is **temporally homogeneous**
- It satisfies the **Markovian property**

That is, for all $t = 0, 1, 2, \dots$ and for every possible set of $t + 1$ states, including S_i and S_j :

$$P(X_{t+1} = j \mid X_0 = k_0, X_1 = k_1, \dots, X_{t-1} = k_{t-1}, X_t = i) = P(X_{t+1} = j \mid X_t = i)$$



Andrey A Markov
Russian Mathematician
1856 - 1922

Markov chains

A sequence of random variables $X_0, X_1, X_2, \dots, X_n, \dots$

A finite state space $\{S_1, \dots, S_k\}$

An initial distribution vector $\pi_0 : P(X_0 = S_i) = \pi_0(i)$

A transition square matrix $T : T_{i,j} = P(X_{n+1} = S_j | X_n = S_i)$

$$\pi_{n+1}(j) = P(X_{n+1} = S_j)$$


$$= \sum_{i=1}^k P(X_n = S_i) P(X_{n+1} = S_j | X_n = S_i)$$

$$= \sum_{i=1}^k \pi_n(i) T_{i,j}$$

$$\begin{bmatrix} T(1,1) & T(1,2) & \cdots & T(1,k) \\ T(2,1) & T(2,2) & & T(2,k) \\ & \vdots & \ddots & \vdots \\ T(k,1) & T(k,2) & \cdots & T(k,k-1) & T(k,k) \end{bmatrix}$$

The transition rule:

$$\pi_{n+1} = \pi_n \cdot T$$



$$T = \begin{array}{c|ccc} & s & c & r \\ \hline s & 0.6 & 0.3 & 0.1 \\ c & 0.2 & 0.3 & 0.5 \\ r & 0.4 & 0.1 & 0.5 \end{array}$$

Independent coin tossing

What is the transition matrix?

Fair coin ($p=0.5$)

$$\begin{pmatrix} 0.5 & 0.5 \\ 0.5 & 0.5 \end{pmatrix}$$

Biased coin ($p=0.7$)

$$\begin{pmatrix} 0.3 & 0.7 \\ 0.3 & 0.7 \end{pmatrix}$$

What is the initial probability distribution?

Fair coin ($p=0.5$)

$$(0.5 \quad 0.5)$$

Biased coin ($p=0.7$)

$$(0.3 \quad 0.7)$$

Transitions further into the future

		Weather on next day		
		Dry	Wet	Total
Weather on one day	Dry	57	12	69
	Wet	12	8	20

$$T = \begin{array}{c|cc} & 0 & 1 \\ \hline 0 & 0.826 & 0.174 \\ 1 & 0.600 & 0.400 \end{array} = \begin{bmatrix} T(0,0) & T(0,1) \\ T(1,0) & T(1,1) \end{bmatrix}$$

Two days into the future:

$$P(X_2 = 0 \mid X_0 = 1) = ?$$

Transitions further into the future

$$P(X_2 = 0 | X_0 = 1) = ?$$

$$\pi_1 = \pi_0 \cdot T = (0,1) \cdot T = (T(1,0), T(1,1))$$

$$P(X_2 = 0 | X_0 = 1) = \pi_2(0) = \pi_1 \cdot \begin{bmatrix} T(0,0) \\ T(1,0) \end{bmatrix} = T(1,0)T(0,0) + T(1,1)T(1,0)$$

And therefore, we get :

$$0.6 \cdot 0.826 + 0.4 \cdot 0.6 = 0.7356$$

$$T = \begin{bmatrix} T(0,0) & T(0,1) \\ T(1,0) & T(1,1) \end{bmatrix}$$

	0	1
0	0.826	0.174
1	0.600	0.400

Transitions further into the future: $T^{(r)} = T^r$

What if we want to talk about day 7?

Note that the above calculation actually means that:

$$\pi_2 = \pi_1 \cdot T = (\pi_0 \cdot T) \cdot T = \pi_0 \cdot T^2$$

and we can continue to get

$$\pi_r = \pi_0 \cdot T^r$$

which we can summarize as: $T^{(r)} = T^r$

The stationary distribution

A probability distribution, σ , over the state space $\{S_1, \dots, S_k\}$ that satisfies:

$$\sigma \cdot T = \sigma$$

$$\pi_0 = \sigma \rightarrow \pi_1 = \pi_0 \cdot T = \sigma \cdot T = \sigma = \pi_0 \rightarrow \pi_n = \sigma \forall n$$


The stationary distribution

$$\sigma \cdot T = \sigma$$

- 1 is an eigenvalue of T
- σ is a left eigenvector of T , with eigenvalue 1


$$\begin{pmatrix} \end{pmatrix} \begin{pmatrix} \end{pmatrix} = 1 \begin{pmatrix} \end{pmatrix}$$
$$\sum_{j=1}^k T_{i,j} v_j = v_i \rightarrow v = \vec{\mathbf{1}}$$

What is the stationary distribution here?



$$T = \begin{bmatrix} 0.3 & 0.7 \\ 0.8 & 0.2 \end{bmatrix}$$

What is the stationary distribution here?



$$T = \begin{bmatrix} 0.3 & 0.7 \\ 0.8 & 0.2 \end{bmatrix}$$

$$(x \quad y)T = (x \quad y) \begin{pmatrix} 0.3 & 0.7 \\ 0.8 & 0.2 \end{pmatrix} = (x \quad y)$$

$$0.3x + 0.8y = x \rightarrow x = \frac{0.8y}{0.7}$$

$$0.7x + 0.2y = y$$

$$x + y = 1 \rightarrow \frac{0.8y}{0.7} + y = 1 \rightarrow 1.5y = 0.7 \rightarrow y = \frac{7}{15}, x = \frac{8}{15}$$

Higher order Markov

$$P(X_{t+1} = j \mid X_0 = k_0, X_1 = k_1, \dots, X_{t-1} = k_{t-1}, X_t = i) = P(X_{t+1} = j \mid X_t = i)$$

Need more terms in the condition.
Will depend not only on the most recent time but on r recent times.

Higher order Markov

$$P(X_{t+1} = j \mid X_0 = k_0, X_1 = k_1, \dots, X_{t-r+1} = k_{t-r+1}, \dots, X_{t-1} = k_{t-1}, X_t = k_t) =$$

$$P(X_{t+1} = j \mid X_{t-r+1} = k_{t-r+1}, \dots, X_{t-1} = k_{t-1}, X_t = k_t)$$



r most recent events

Markov chains application: text prediction

- Use the previous $N-1$ words in a sequence to predict the next word
- In auto completion and in auto translation these are called N-Grams
- Language Model (LM)
 - unigrams, bigrams, trigrams,...

Using N-Grams

- For N-gram models $P(w_n | w_1^{n-1}) \approx P(w_n | w_{n-N+1}^{n-1})$
- Bi-grams: $N = 2$. A Markov assumption.
- By the Chain Rule we can compute probabilities of sentences:

$$\begin{aligned} P(w_1, w_2, w_3, \dots, w_n) &= \\ &= P(w_n | w_1, w_2, \dots, w_{n-1}) P(w_{n-1} | w_1, w_2, \dots, w_{n-2}) \dots P(w_2 | w_1) P(w_1) \end{aligned}$$

What would this be under a Markov assumption?

$$= P(w_n | w_{n-1}) P(w_{n-1} | w_{n-2}) \dots P(w_2 | w_1) P(w_1)$$

Bigram Counts - example

- Out of 9222 sentences

	i	want	to	eat	chinese	food	lunch	spend
i	5	827	0	9	0	0	0	2
want	2	0	608	1	6	6	5	1
to	2	0	4	686	2	0	6	211
eat	0	0	2	0	16	2	42	0
chinese	1	0	0	0	0	82	1	0
food	15	0	15	0	1	4	0	0
lunch	2	0	0	0	0	1	0	0
spend	1	0	1	0	0	0	0	0

i	want	to	eat	chinese	food	lunch	spend
2533	927	2417	746	158	1093	341	278

Using the Bigram

$$P(\text{"i want chinese food"}) =$$

$$P(i) * P(\text{want}|i) * P(\text{chinese} | \text{want}) * P(\text{food} | \text{chinese}) =$$

$$\frac{\#(i)}{\#(\text{all sentences})} * \frac{\#(\text{"i want"})}{\#(i)} * \frac{\#(\text{"want chinese"})}{\#(\text{want})} * \frac{\#(\text{"chinese food"})}{\#(\text{chinese})} =$$

$$\left(\frac{2533}{9222}\right) * \left(\frac{827}{2533}\right) * \left(\frac{6}{927}\right) * \left(\frac{82}{158}\right) = 0.0003$$

Useful for ...

- Speech recognition:
“I ate a cherry” is a more likely sentence than “Eye eight a Jerry”
- Machine translation
 $\text{Pr}[\text{high acceptance rate}] > \text{Pr}[\text{tall acceptance rate}]$
 $\text{Pr}[\text{finite list}] >? \text{Pr}[\text{final list}]$
- Context sensitive spelling correction
“Their are problems wit this sentence.”
- Sentence completion “Please turn off your ...” , “I want to eat ...”

Summary

- Brief intro to Markov Chains
- The transition matrix and its powers
- The stationary distribution
- N-grams
- Covariance and the CLT